

VIRTUAL MEMORY

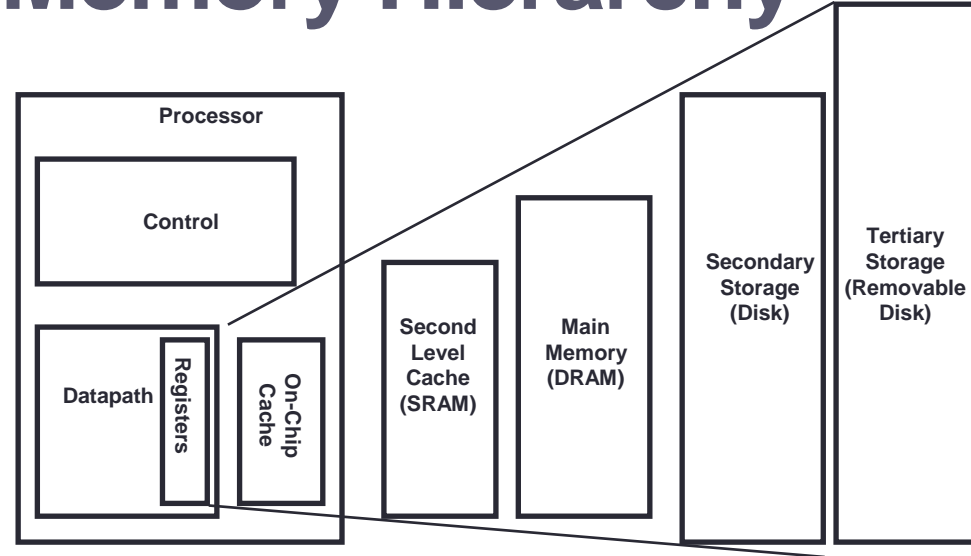
Lecturer: Hui Annie Guo

h.guo@unsw.edu.au

K17-501F

Recall: Overview of Memory Hierarchy

- **registers \leftrightarrow memory**
 - by compiler/programmer
- **cache \leftrightarrow memory**
 - by the hardware
- **memory \leftrightarrow disks**
 - by the hardware and operating system (virtual memory)
 - by the programmer



Lecture overview

- **Topics**


- A glance of virtual memory
- A hardware-centric view
 - Page table
 - TLB

- **Suggested reading**

- H&P Chapter 5.7
- https://en.wikipedia.org/wiki/Virtual_memory

Virtual memory

- **A memory management technique**
 - Use the main memory and disk to create an illusion of very large memory to the user.
- **Use main memory as a “cache” for the secondary (disk) storage**
 - Managed mainly by the operating system (OS)



Virtual memory (cont.)

- Programs have separate **virtual memory space**
- A process is a program that is being executed
- Processes share the main memory
 - Each gets a private address space, holding its code and data
 - They are protected from each other by OS

Virtual memory (cont.)

- **Page**
 - A virtual block
- **Page hit**
 - Data accessed are in the main memory
- **Page fault**
 - Data accessed are not in the main memory

Virtual memory (cont.)

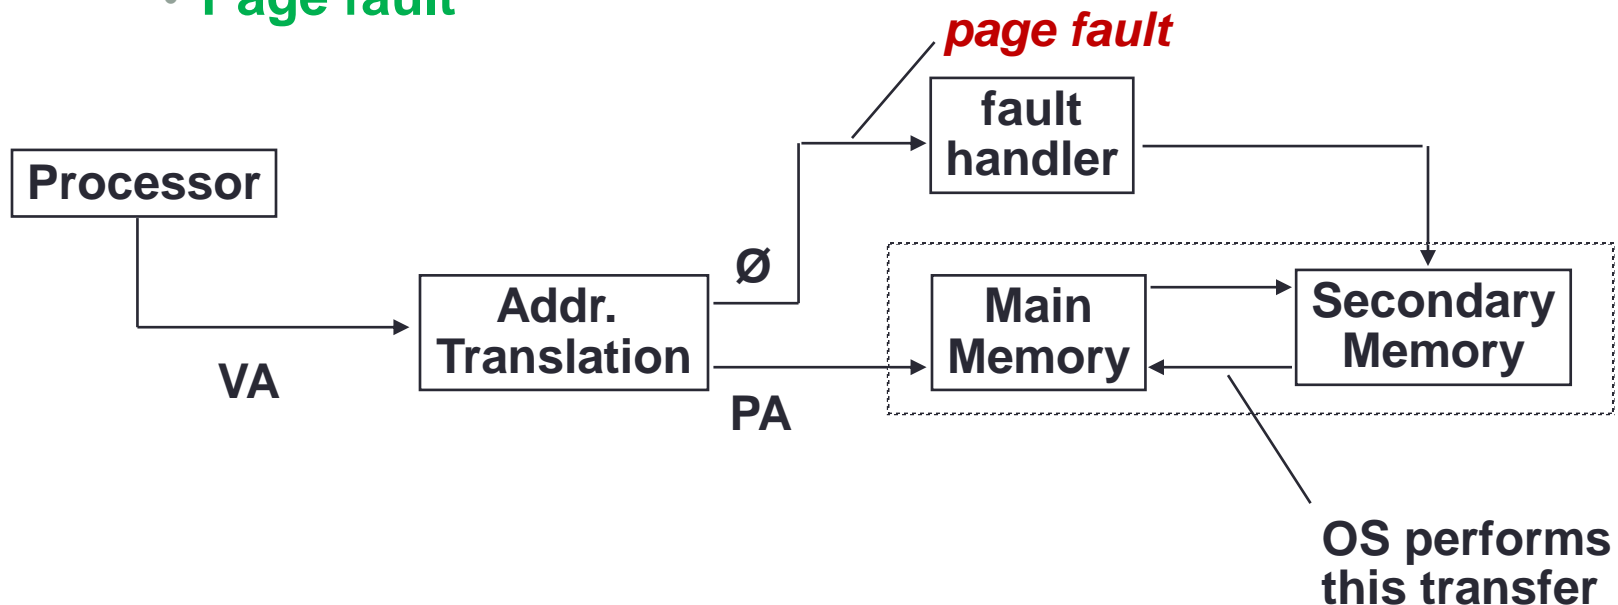
- The basic tasks of using virtual memory
 - Translation between **virtual address (VA)** and **physical address (PA)**
 - Access control of physical memory space
- The control of using virtual memory is often carried out in a component, called memory management unit (MMU)
 - VM translation “miss” is effectively a page fault

Four design issues

- **Where to place a page?**
 - Fully associative or highly associative
- **How to find a page?**
 - Address translation
- **Which page should be replaced on a page fault?**
 - Sophisticated (LRU + “working set”)
- **What happens when you want to write a page back?**
 - Always write-back and write allocate
 - Disk writes take millions of clock cycles

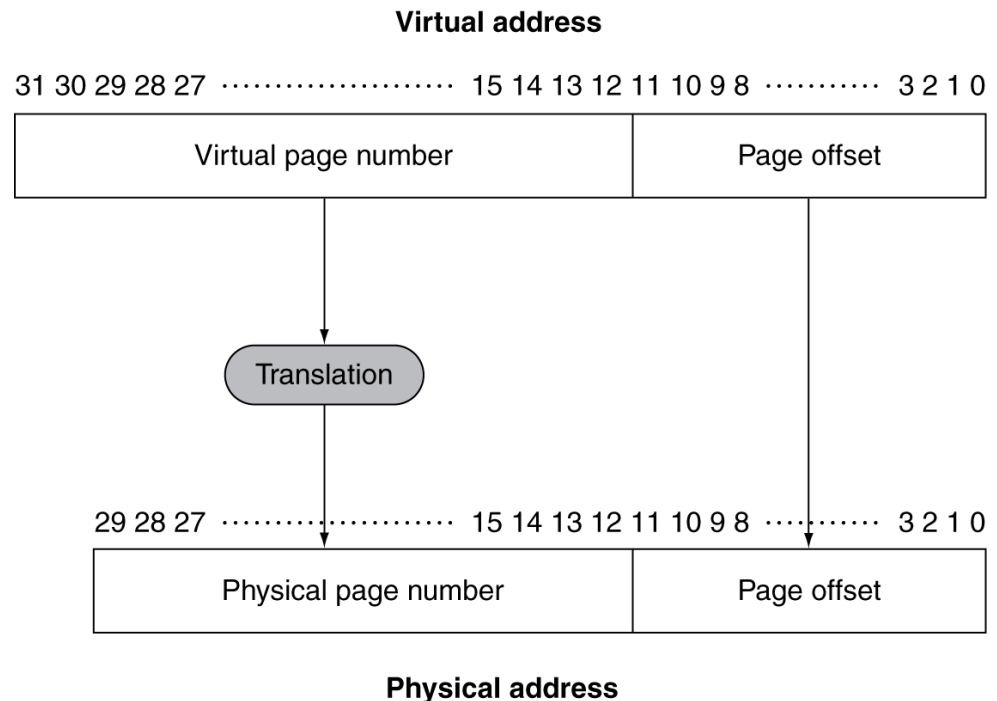

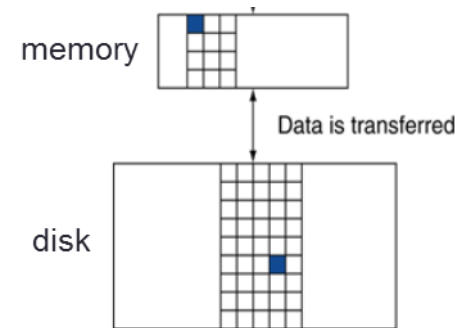
Memory access with virtual address

- **Given a virtual address, the address translation generates two possible results**
 - **PA**
 - **Page hit**
 - **no PA**
 - **Page fault**



Address translation


- **Virtual address**
 - Virtual page number + Page offset
- **Physical address**
 - Physical page number + Page offset
- **Translation is based on page table**



Page table (PT)

- **Usually is implemented in the physical memory.**
- **Stores page placement information**
 - An array of page table entries (PTE), indexed by virtual page number (VPN)
 - Page table register (PTR) points to the page table location in the memory
- **If a page is present in memory**
 - PTE stores the physical page number
 - Plus other status bits (referenced, dirty, ...)
- **If a page is not present**
 - PTE refers to the location in the swap space on disk

Translation using page table



In-class exercise (1)

- **Assume a program is executed in a virtual memory space with the page size of 8KB without page fault. The PTR for the execution is set to 0x00110. The figure to the right shows the contents of a set of main memory locations.**
 - **(a) What is the Page offset field size in the virtual address?**

| main memory | |
|-------------|----------|
| ADDR | CONTENTS |
| 0x00000 | 0x0000 |
| 0x00100 | 0x0010 |
| 0x00110 | 0x0022 |
| 0x00120 | 0x0045 |
| 0x00130 | 0x0078 |
| 0x00145 | 0x0010 |
| 0x10000 | 0x2333 |
| 0x10020 | 0x4444 |
| 0x22000 | 0x1111 |
| 0x22020 | 0x2222 |
| 0x45000 | 0x5555 |
| 0x45020 | 0x6676 |

In-class exercise (1)

- **Assume a program is executed in a virtual memory space with the page size of 8KB without page fault. The PTR for the execution is set to 0x00110. Figure to the right shows the contents of a set of main memory locations.**
 - **(b) Does a PA exist for the following virtual address?**
0x010020

| main memory | |
|-------------|----------|
| ADDR | CONTENTS |
| 0x00000 | 0x0000 |
| 0x00100 | 0x0010 |
| 0x00110 | 0x0022 |
| 0x00120 | 0x0045 |
| 0x00130 | 0x0078 |
| 0x00145 | 0x0010 |
| 0x10000 | 0x2333 |
| 0x10020 | 0x4444 |
| 0x22000 | 0x1111 |
| 0x22020 | 0x2222 |
| 0x45000 | 0x5555 |
| 0x45020 | 0x6676 |

In-class exercise (1)

- **Assume a program is executed in a virtual memory space with the page size of 8KB without page fault. The PTR for the execution is set to 0x00110. Figure to the right shows the contents of a set of main memory locations.**
 - **(c) What value will be returned to processor if it requests to read the following address?**

0x010020

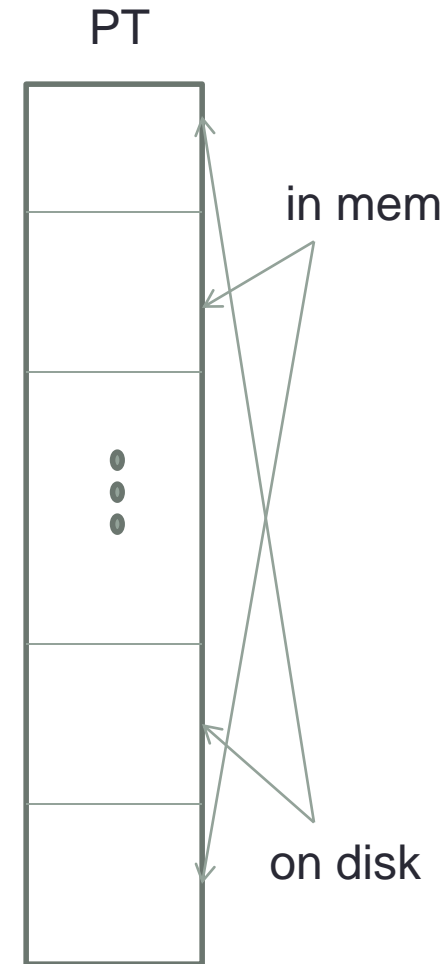
| main memory | |
|-------------|----------|
| ADDR | CONTENTS |
| 0x00000 | 0x0000 |
| 0x00100 | 0x0010 |
| 0x00110 | 0x0022 |
| 0x00120 | 0x0045 |
| 0x00130 | 0x0078 |
| 0x00145 | 0x0010 |
| 0x10000 | 0x2333 |
| 0x10020 | 0x4444 |
| 0x22000 | 0x1111 |
| 0x22020 | 0x2222 |
| 0x45000 | 0x5555 |
| 0x45020 | 0x6676 |

In-class exercise (2)

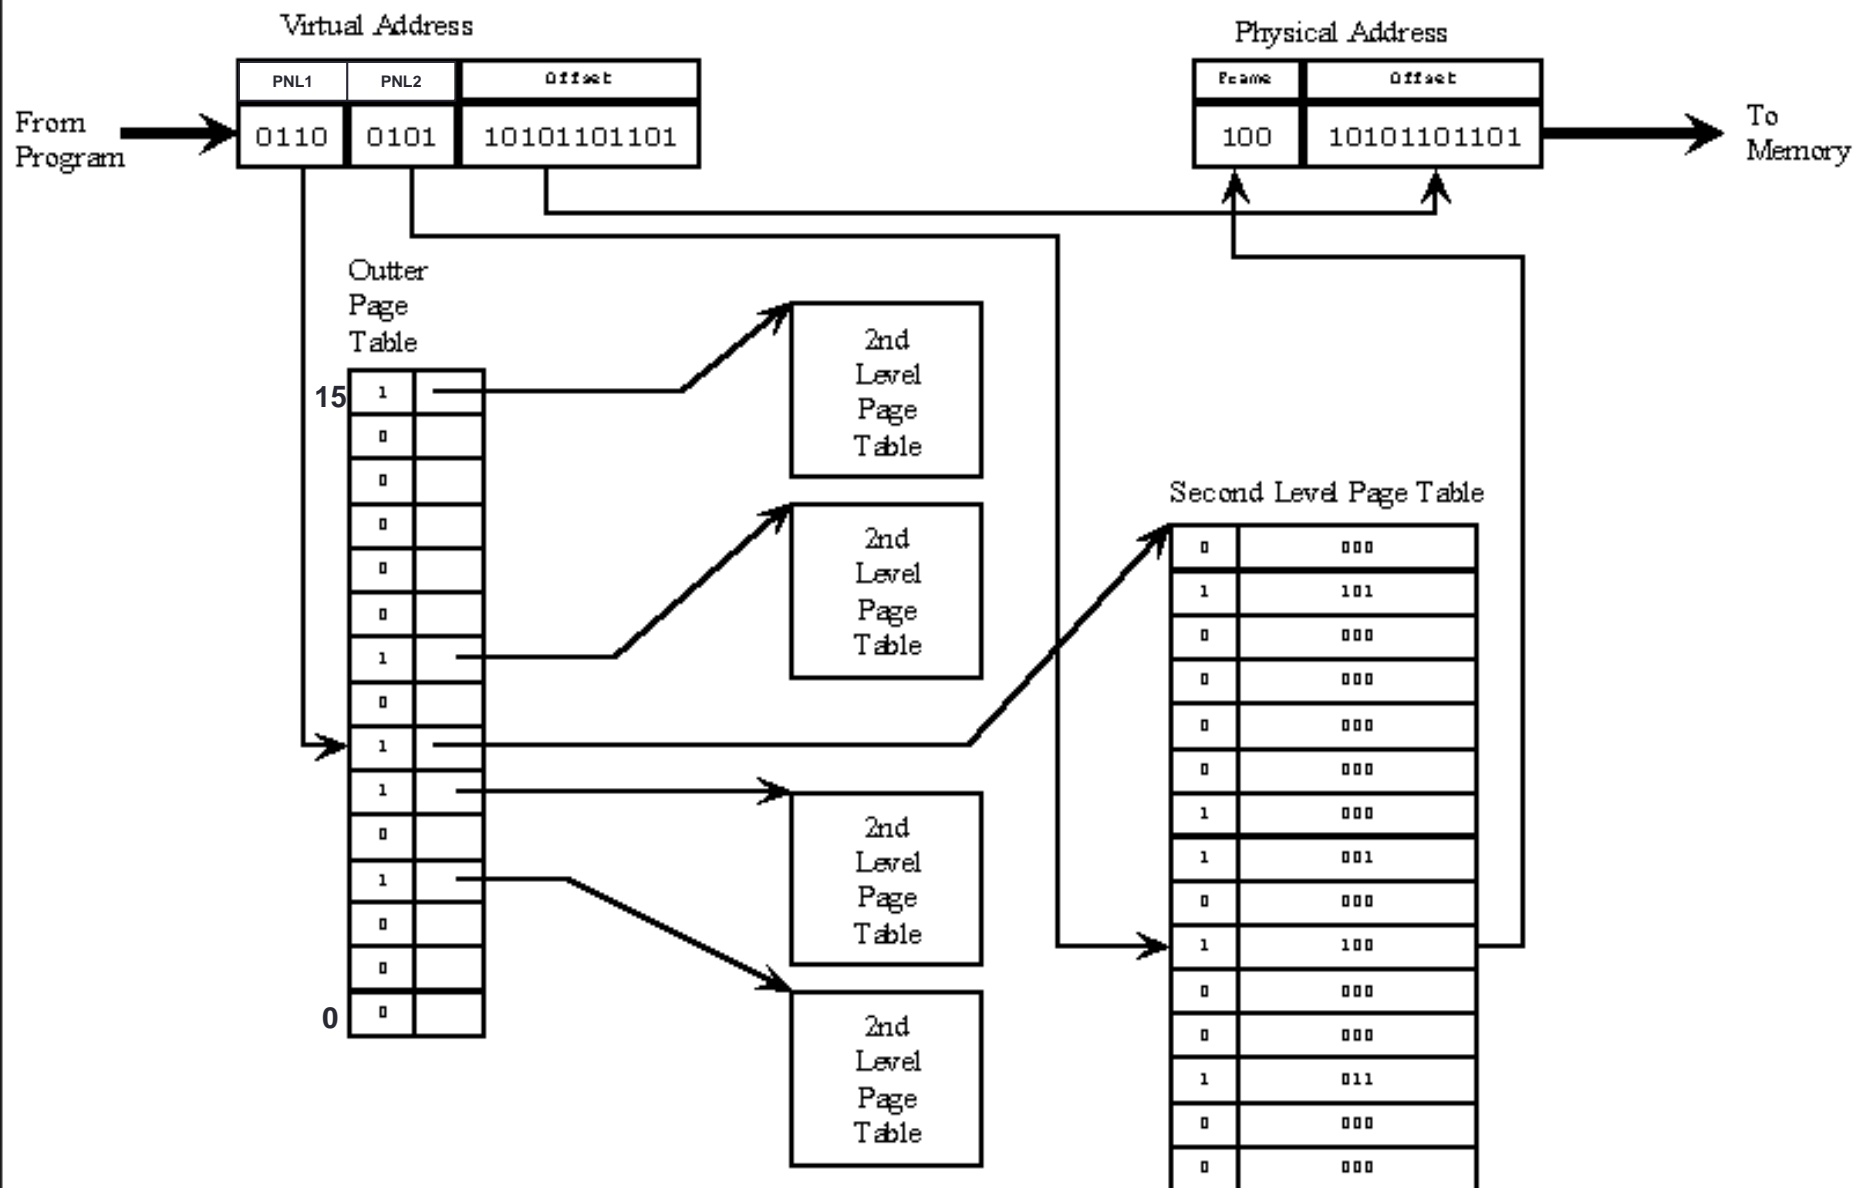
- **For the virtual memory space, physical memory size and the page size given in exercise (1), what is the size of a page table?**
- **What if the virtual address is increased to 32 bits?**

Multiple level page table*

- **Only store a fraction of the table in the memory**
 - **Dividing the page table into a set of sub-tables**
 - **Some sub-tables stored in the memory, and**
 - **Other sub-tables paged out to the disk**



Translation with multi-level page table*



Fast translation using TLB

- **Address translation incurs extra memory references**
 - **access PT**
 - **access the actual memory**
- **But access to page table has good locality**
 - **Leading to the TLB design**

TLB: Translation look-aside buffer

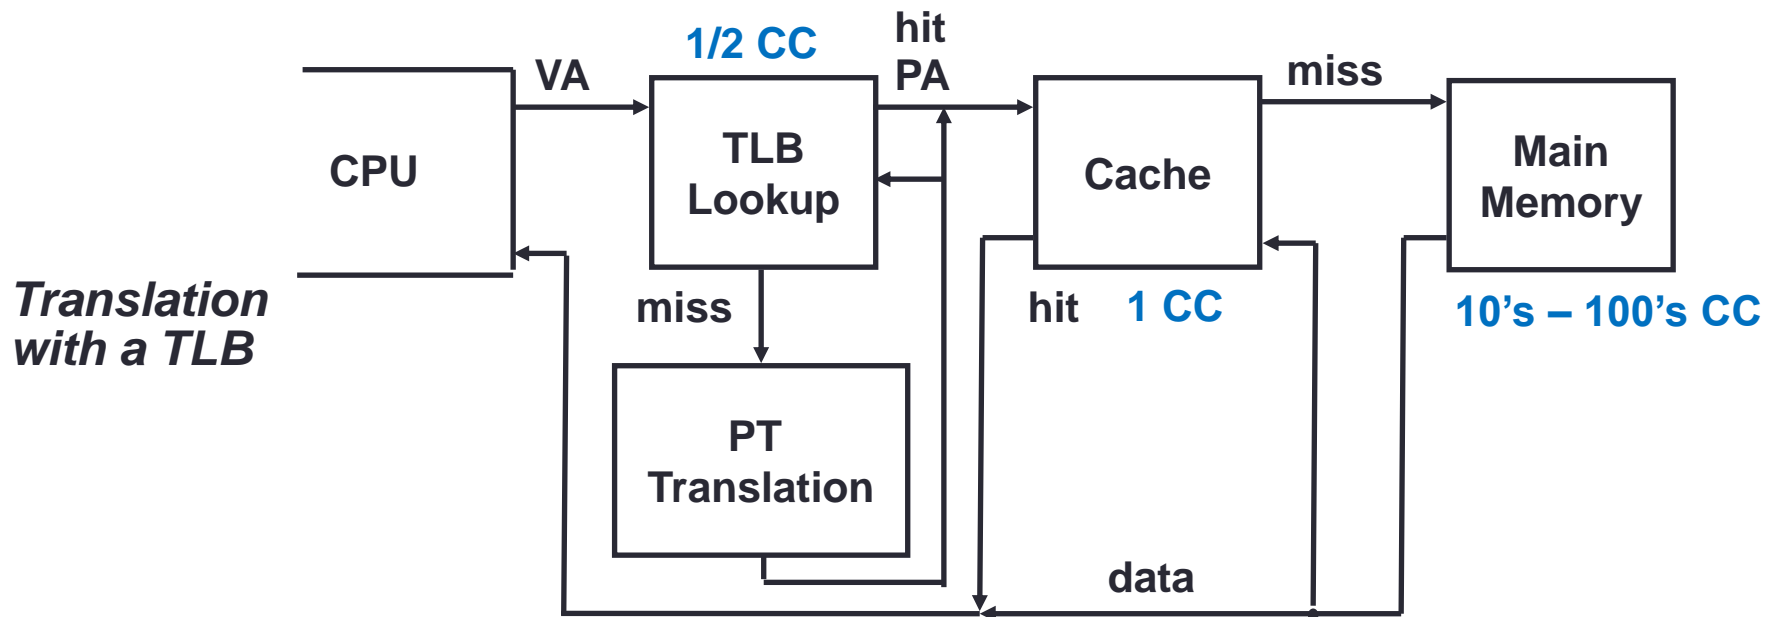
- **Is a special cache for page table**
- **Implemented on the processor chip**
- **Offers a fast VA to PA translation**
- **Just like any other cache, TLB can be organized as**
 - **fully associative,**
 - **set associative, or**
 - **direct mapped**

TLB: Translation look-aside buffer (cont.)

- **TLBs are usually small, typically not more than 128~256 entries**
 - **This permits fully associative lookup on high-end machines.**
 - **Most mid-range machines use small n-way set associative organizations.**

CPU datapath with TLB

- The delay of TLB lookup is less than one clock cycle
 - E.g. 1/2 clock cycles, as compared to many clock cycles of the page-table based translation



Further improvement

- **Reduce the impact of address translation on performance**
 - Use the slack time in a pipeline stage execution
 - Overlap TLB access with another processor operation

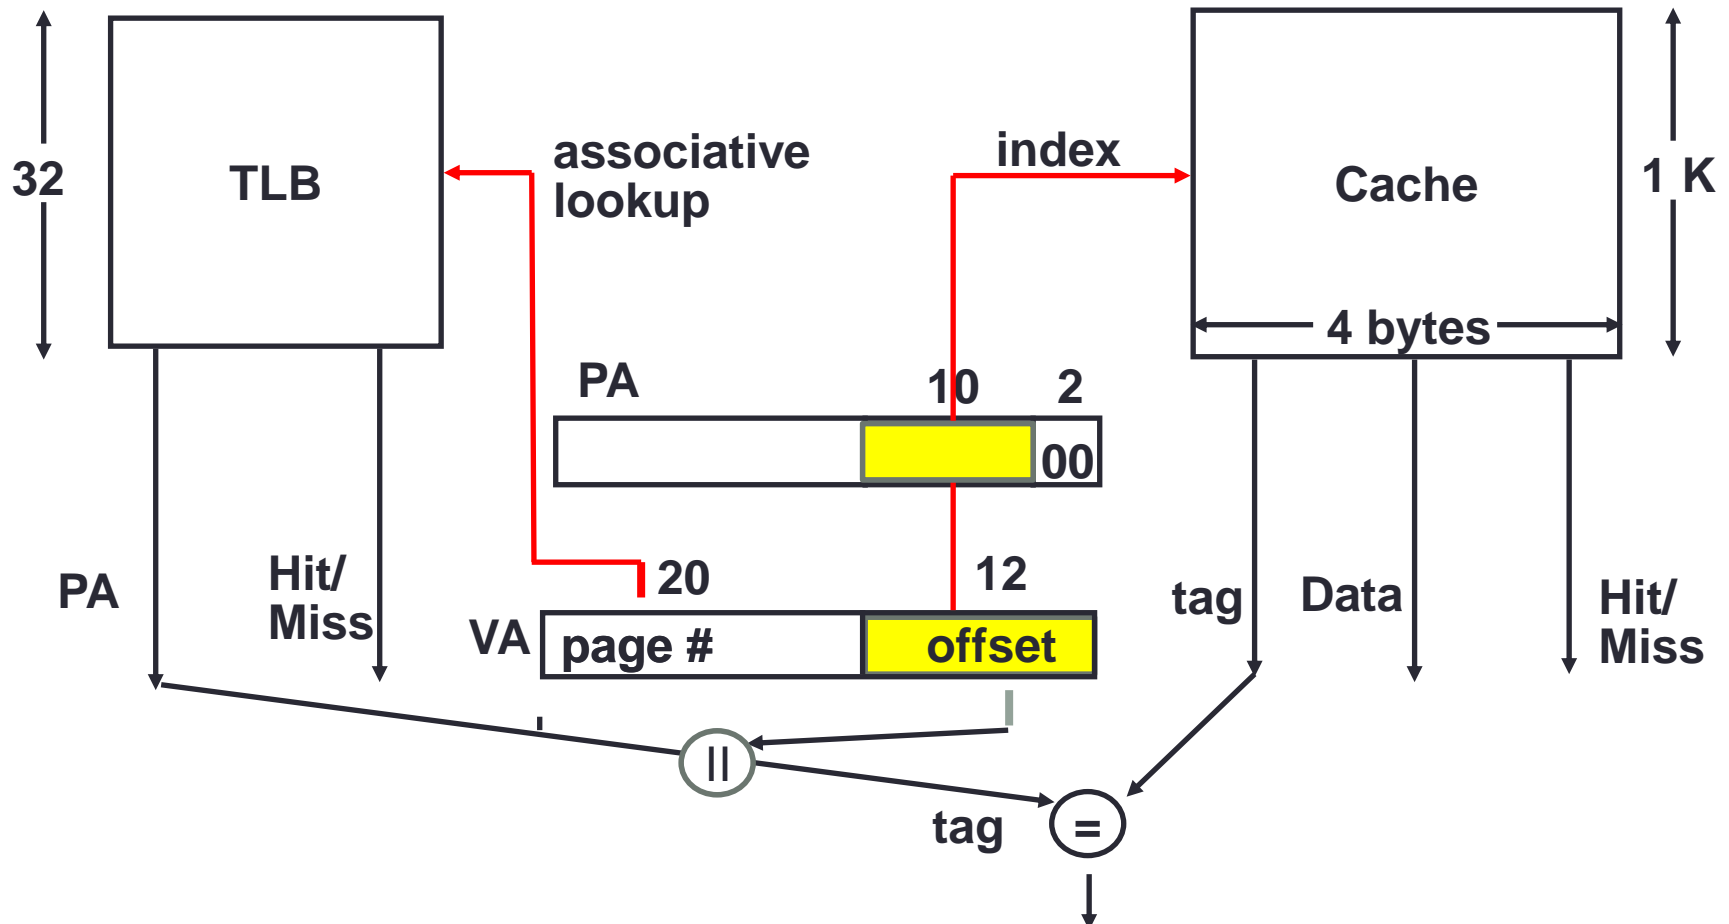
MIPS R3000 Pipeline

| Inst Fetch | | Dcd/ Reg | | ALU / V.A | | Memory | Write Reg | |
|------------|--|----------|--|-----------|-----|---------|-----------|--|
| TLB | | RF | | Operation | | | WB | |
| I-Cache | | | | V.A. | TLB | D-Cache | | |

Overlap I-cache with TLB access

- **High order bits of VA are used to look up in TLB**
- **Low order bits of VA are used as index to search in cache**
- **Both can be performed in parallel**
- **See next slide for demonstration**

Overlap I-cache with TLB access (cont.)



IF cache hit AND (cache tag = PA) THEN deliver data to CPU
 ELSE IF cache miss AND TLB hit THEN
 access memory with the PA from the TLB
 ELSE do standard VA translation

Summary of memory hierarchy

- **Common principles apply at all levels of the memory hierarchy**
 - **Based on the notion of caching**
- **At each level in the hierarchy, four issues should be addressed**
 - **Block placement**
 - **Block identification**
 - **Replacement on a miss**
 - **Write policy**